

# M5: Outliers and Missing Data

Luana Lima

02/11/2021

## Setting R code chunk options

First R code chunk is used for setting the options for all R code chunks. The choice `echo=TRUE` means both code and output will appear on report, `include = FALSE` neither code nor output is printed.

## Loading packages and initializing

Second R code chunk is for loading packages. By setting `message = FALSE`, the code will appear but not the output.

```
library(lubridate)
library(ggplot2)
library(forecast)
library(Kendall)
library(tseries)

#New packages for M5
#install.packages("outliers")
library(outliers)
#install.packages("tidyverse")
library(tidyverse)
```

## Importing data

Let's continue working with our inflow data for reservoirs in Brazil.

```
#Importing time series data from text file#
raw_inflow_data <- read.table(file="../Data/inflowtimeseries.txt",header=FALSE,skip=0)

#Trim the table to include only columns you need
nhydro <- ncol(raw_inflow_data)-2
nobs <- nrow(raw_inflow_data)

#If your file does not have header like this one you can add column names after
#creating the data frame
colnames(raw_inflow_data)=c("Month", "Year", "HP1", "HP2", "HP3", "HP4", "HP5",
                             "HP6", "HP7", "HP8", "HP9", "HP10", "HP11", "HP12",
                             "HP13", "HP14", "HP15")

#Checking data
head(raw_inflow_data)
```

```
##   Month Year  HP1  HP2  HP3  HP4  HP5  HP6 HP7  HP8 HP9 HP10 HP11 HP12 HP13
## 1   Jan 1931 4782 4076 2518 2450 2649 1462 450  968 246 2636  452 4870  452
```

```
## 2 Feb 1931 7323 7681 4188 150 2401 758 554 219 74 4158 457 4550 796
## 3 Mar 1931 8266 5921 3253 2389 3261 707 615 333 123 3847 631 6537 804
## 4 Apr 1931 6247 4600 2449 1253 2006 469 474 297 113 3291 510 7298 644
## 5 May 1931 3642 2789 1651 2374 2454 3167 378 3295 938 1956 276 4942 421
## 6 Jun 1931 2425 2062 1270 2672 2433 3236 301 2547 951 1371 201 2478 305
## HP14 HP15
## 1 17342 31270
## 2 21530 43827
## 3 33299 49884
## 4 34674 43962
## 5 15184 35156
## 6 8611 25764
```

```
str(raw_inflow_data)
```

```
## 'data.frame': 972 obs. of 17 variables:
## $ Month: chr "Jan" "Feb" "Mar" "Apr" ...
## $ Year : int 1931 1931 1931 1931 1931 1931 1931 1931 1931 1931 ...
## $ HP1 : int 4782 7323 8266 6247 3642 2425 2158 1854 1839 1896 ...
## $ HP2 : int 4076 7681 5921 4600 2789 2062 1644 1301 1439 1340 ...
## $ HP3 : int 2518 4188 3253 2449 1651 1270 1204 1152 1297 1259 ...
## $ HP4 : int 2450 150 2389 1253 2374 2672 1238 605 1016 674 ...
## $ HP5 : int 2649 2401 3261 2006 2454 2433 1798 1160 1584 1563 ...
## $ HP6 : int 1462 758 707 469 3167 3236 1957 844 1937 1484 ...
## $ HP7 : int 450 554 615 474 378 301 256 244 222 355 ...
## $ HP8 : int 968 219 333 297 3295 2547 2585 1173 3596 1140 ...
## $ HP9 : int 246 74 123 113 938 951 883 404 378 211 ...
## $ HP10 : int 2636 4158 3847 3291 1956 1371 1186 1049 1162 1507 ...
## $ HP11 : int 452 457 631 510 276 201 213 196 161 208 ...
## $ HP12 : int 4870 4550 6537 7298 4942 2478 1905 1647 1453 1358 ...
## $ HP13 : int 452 796 804 644 421 305 261 246 250 328 ...
## $ HP14 : int 17342 21530 33299 34674 15184 8611 5939 4259 3282 3305 ...
## $ HP15 : int 31270 43827 49884 43962 35156 25764 18109 13320 8225 8900 ...
```

## Creating the date object

Here we use the function `my()` from package `lubridate`.

```
#using package lubridate
```

```
my_date <- paste(raw_inflow_data[,1],raw_inflow_data[,2],sep="-")
my_date <- my(my_date) #function my from package lubridate
head(my_date)
```

```
## [1] "1931-01-01" "1931-02-01" "1931-03-01" "1931-04-01" "1931-05-01"
## [6] "1931-06-01"
```

```
#add that to inflow_data and store in a new data frame
```

```
inflow_data <- cbind(my_date,raw_inflow_data[,3:(3+nhydro-1)])
head(inflow_data)
```

```
## my_date HP1 HP2 HP3 HP4 HP5 HP6 HP7 HP8 HP9 HP10 HP11 HP12 HP13
## 1 1931-01-01 4782 4076 2518 2450 2649 1462 450 968 246 2636 452 4870 452
## 2 1931-02-01 7323 7681 4188 150 2401 758 554 219 74 4158 457 4550 796
## 3 1931-03-01 8266 5921 3253 2389 3261 707 615 333 123 3847 631 6537 804
## 4 1931-04-01 6247 4600 2449 1253 2006 469 474 297 113 3291 510 7298 644
## 5 1931-05-01 3642 2789 1651 2374 2454 3167 378 3295 938 1956 276 4942 421
## 6 1931-06-01 2425 2062 1270 2672 2433 3236 301 2547 951 1371 201 2478 305
```

```
##      HP14  HP15
## 1 17342 31270
## 2 21530 43827
## 3 33299 49884
## 4 34674 43962
## 5 15184 35156
## 6  8611 25764
```

## Removing zeros in the end on data

```
#Remove last for rows by replacing current data frame
inflow_data <- inflow_data[1:(nobs-4),]
my_date <- my_date[1:(nobs-4)]

#update object with number of observations
nobs <- nobs-4

#Tail again to check if the rows were correctly removed
tail(inflow_data)
```

```
##      my_date  HP1  HP2  HP3  HP4  HP5  HP6  HP7  HP8  HP9  HP10  HP11  HP12  HP13
## 963 2011-03-01 8897 5426 5805 2009 3576 1834 798 2097 1071 3435 797 3693 943
## 964 2011-04-01 4991 3207 3323 4063 3235 1620 481 2325 902 2173 493 5255 563
## 965 2011-05-01 3025 2156 2274 2351 2063 572 304 1496 540 1175 254 1998 415
## 966 2011-06-01 2415 1813 1936 1836 2087 713 270 2294 898 985 130 1256 311
## 967 2011-07-01 1883 1426 1560 2930 2105 2988 233 4578 2045 864 119 1068 275
## 968 2011-08-01 1444 1139 1441 5069 2328 4559 224 4573 2527 827 120 854 251
##      HP14  HP15
## 963 29976 39843
## 964 28892 39441
## 965 20978 31023
## 966  7081 21840
## 967  3910 14162
## 968  2561  8896
```

## Transforming data into time series object

Many of the functions we will use require a time series object. You can transform your data in a time series using the function `ts()`.

```
ts_inflow_data <- ts(inflow_data[,2:(2+nhydro-1)],frequency=12)
#note that we are only transforming columns with inflow data, not the date columns #start=my_date[1],e
head(ts_inflow_data,15)
```

```
##      HP1  HP2  HP3  HP4  HP5  HP6  HP7  HP8  HP9  HP10  HP11  HP12  HP13  HP14
## Jan 1 4782 4076 2518 2450 2649 1462 450 968 246 2636 452 4870 452 17342
## Feb 1 7323 7681 4188 150 2401 758 554 219 74 4158 457 4550 796 21530
## Mar 1 8266 5921 3253 2389 3261 707 615 333 123 3847 631 6537 804 33299
## Apr 1 6247 4600 2449 1253 2006 469 474 297 113 3291 510 7298 644 34674
## May 1 3642 2789 1651 2374 2454 3167 378 3295 938 1956 276 4942 421 15184
## Jun 1 2425 2062 1270 2672 2433 3236 301 2547 951 1371 201 2478 305 8611
## Jul 1 2158 1644 1204 1238 1798 1957 256 2585 883 1186 213 1905 261 5939
## Aug 1 1854 1301 1152 605 1160 844 244 1173 404 1049 196 1647 246 4259
## Sep 1 1839 1439 1297 1016 1584 1937 222 3596 378 1162 161 1453 250 3282
## Oct 1 1896 1340 1259 674 1563 1484 355 1140 211 1507 208 1358 328 3305
```


```
## Nov 1 2095 1447 1218 674 1404 835 371 563 252 1996 596 1905 319 6500
## Dec 1 2725 2479 2013 1278 2272 1073 419 512 197 3015 381 2121 335 8461
## Jan 2 4679 4021 2435 1259 1995 1044 520 609 159 3978 711 3811 467 14002
## Feb 2 5535 4082 2262 1895 2996 1454 525 1219 268 2615 316 4681 531 20596
## Mar 2 4310 3398 2065 1686 2392 1888 674 1332 304 2269 271 3329 501 21638
##      HP15
## Jan 1 31270
## Feb 1 43827
## Mar 1 49884
## Apr 1 43962
## May 1 35156
## Jun 1 25764
## Jul 1 18109
## Aug 1 13320
## Sep 1 8225
## Oct 1 8900
## Nov 1 13766
## Dec 1 20880
## Jan 2 33160
## Feb 2 39791
## Mar 2 48274
```

## Initial plots for outlier detection


Common plots for outlier detection are histograms and boxplots. Histograms will help you understand the shape and spread of the data and to identify any potential outliers. And boxplots will give more information on the spread of the data.

```
#using package ggplot2 to make histograms
for(i in 1:nhydro){
  print(ggplot(inflow_data, aes(inflow_data[, (1+i)])) +
        geom_histogram() +
        xlab(paste0("Inflow ", colnames(inflow_data)[(1+i)], sep=""))
  )
}


## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```







```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



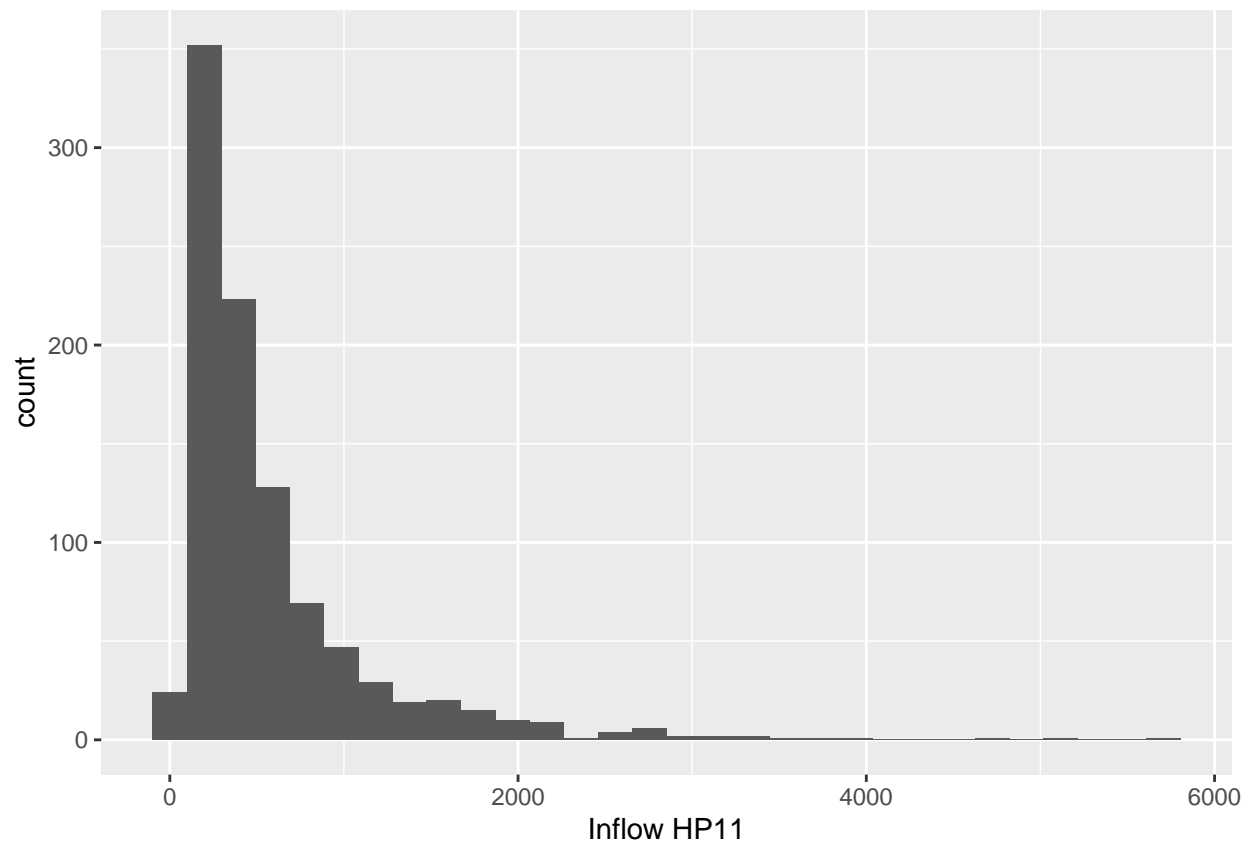
```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```








```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```







```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

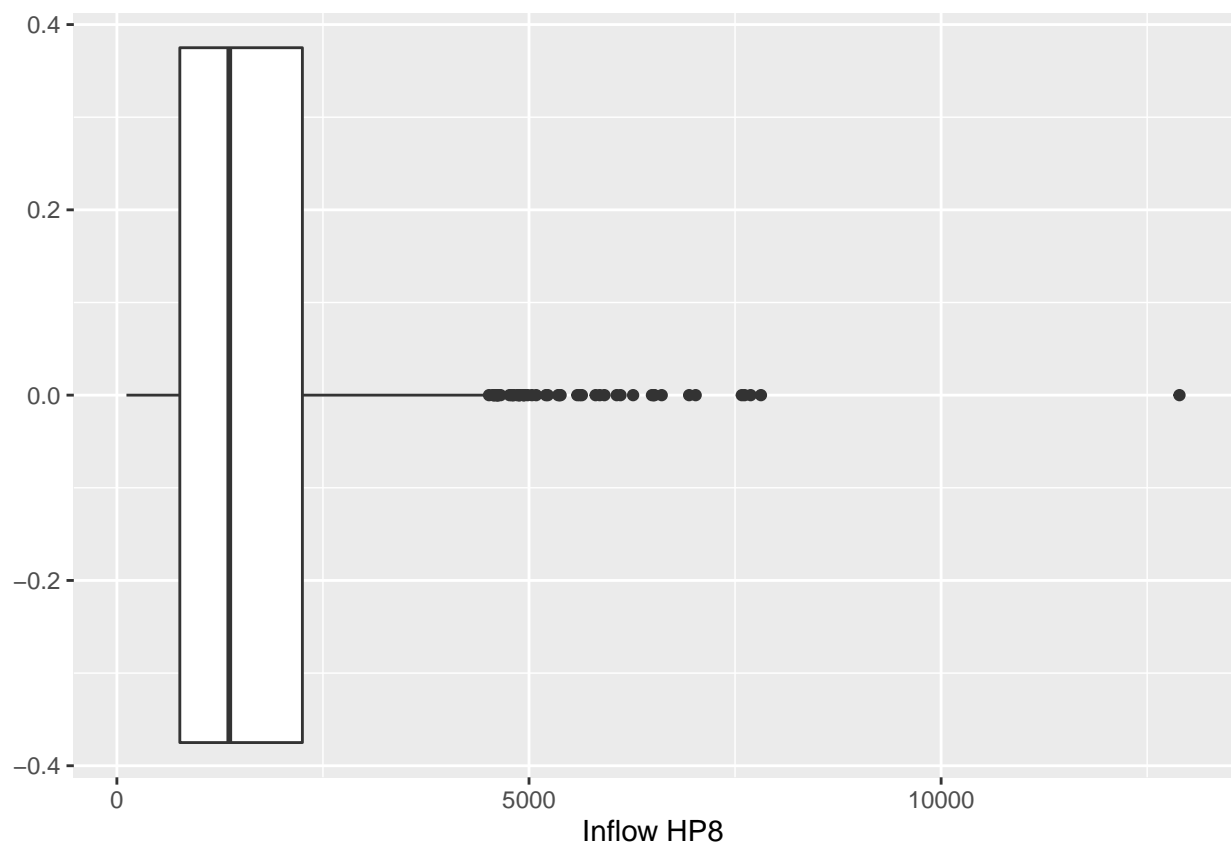





```
#using package ggplot2 to make boxplots
for(i in 1:nhydro){
  print(ggplot(inflow_data, aes(inflow_data[, (1+i)])) +
    geom_boxplot() +
    xlab(paste0("Inflow ", colnames(inflow_data)[(1+i)], sep=" "))
  )
}
```
























Note from the plots that some reservoirs have many points outside the box. But it's hard to tell if those are outliers or not because we are looking at the time series with all its components.


## Decomposing the time series

The stats package has a function called `decompose()`. This function only take time series object. As the name says the decompose function will decompose your time series into three components: trend, seasonal and random. This is similar to what we did in the previous script, but in a more automated way.

The random component is the time series without seasonal and trend component. Let's try to identify outliers by looking at the random component only.

```
#Using R decompose function
iHP=1
decompose_inflow_data=decompose(ts_inflow_data[,iHP], "additive")
plot(decompose_inflow_data)
```

## Decomposition of additive time series



```
#Inspect random component
inflow_random <- decompose_inflow_data$random
mean_inflow <- mean(inflow_random)
sd_inflow <- sd(inflow_random)

cat(mean_inflow,sd_inflow)
```

```
## NA NA
```

```
#Note random series has some missing values, that is why we got NAs
```

```
#Compute mean and standard deviation without missing values
```

```
mean_inflow <- mean(na.exclude(inflow_random)) #exclude NA or missing observation to compute mean and
sd_inflow <- sd(na.exclude(inflow_random))
```

```
cat(mean_inflow,sd_inflow)
```

```
## -4.839207 764.0217
```

## Missing observations

The decompose function introduced NAs in the beginning and end of the data set. Let's just remove them. NAs on the tails can be simply removed.

```
#Create data frame for further use with new random series
```

```
inflow_random <- data.frame(date=my_date,month=as.factor(month(my_date)),inflow=as.numeric(inflow_random))
```

```
#How many NAs we have, you can get it from summary or using is.na()
```

```
sum(is.na(inflow_random$inflow))
```

```
## [1] 12
```

```
#We have NAs in the beginning and end of data, just remove them
head(inflow_random,10)
```

```
##           date month      inflow
## 1  1931-01-01     1          NA
## 2  1931-02-01     2          NA
## 3  1931-03-01     3          NA
## 4  1931-04-01     4          NA
## 5  1931-05-01     5          NA
## 6  1931-06-01     6          NA
## 7  1931-07-01     7 -382.56004
## 8  1931-08-01     8 -340.76212
## 9  1931-09-01     9  -28.44389
## 10 1931-10-01    10  100.05975
```

```
tail(inflow_random,10)
```

```
##           date month      inflow
## 959 2010-11-01    11   53.03006
## 960 2010-12-01    12 -607.79650
## 961 2011-01-01     1  380.88423
## 962 2011-02-01     2 -1857.77671
## 963 2011-03-01     3          NA
## 964 2011-04-01     4          NA
## 965 2011-05-01     5          NA
## 966 2011-06-01     6          NA
## 967 2011-07-01     7          NA
## 968 2011-08-01     8          NA
```

```
#Just remove them
inflow_random <- na.omit(inflow_random)
```

```
#Check data again
sum(is.na(inflow_random$inflow))
```

```
## [1] 0
```

```
head(inflow_random,10)
```

```
##           date month      inflow
## 7  1931-07-01     7 -382.56004
## 8  1931-08-01     8 -340.76212
## 9  1931-09-01     9  -28.44389
## 10 1931-10-01    10  100.05975
## 11 1931-11-01    11 -189.76160
## 12 1931-12-01    12 -813.67150
## 13 1932-01-01     1  -29.36577
## 14 1932-02-01     2  660.34829
## 15 1932-03-01     3 -430.83024
## 16 1932-04-01     4 -552.04754
```

```
tail(inflow_random,10)
```


```
##           date month      inflow
## 953 2010-05-01     5  -304.991106
## 954 2010-06-01     6    3.317439
## 955 2010-07-01     7   25.773293
```

```
## 956 2010-08-01      8    26.029543
## 957 2010-09-01      9   -387.860561
## 958 2010-10-01     10  -378.023582
## 959 2010-11-01     11    53.030064
## 960 2010-12-01     12  -607.796499
## 961 2011-01-01      1   380.884231
## 962 2011-02-01      2 -1857.776707
```


Data is ready!

## Visualizing outliers in R


```
#Generating a box plot by factor where factor is month of the year
ggplot(inflow_random, aes(x=month, y=inflow)) +
  geom_boxplot()
```



```
ggplot(inflow_random, aes(x=date, y=inflow)) +
  geom_line() +
  geom_abline(slope=0, intercept=3*sd_inflow, color="red") +
  geom_abline(slope=0, intercept=-3*sd_inflow, color="red")
```



```
ggplot(inflow_random, aes(y=inflow)) +  
  geom_boxplot()
```



Since we removed the seasonal and trend component, the mean of the random series should be close to zero. Note that from the line plot with the red lines we see that we do have some outliers. The outliers could be due to error collecting the data or an extreme event. Either way, we may want to remove/replace them before fitting a model to our data set to avoid the effect of outliers on our model coefficients.

The box plots are showing more detailed information about the probability distribution for each month of the year. Note that the same months have larger standard deviations.

## Using pre-built functions for outlier detection

We will explore a few function for outlier detection in R.

`outlier()`: this function identifies the value that deviates the most from the mean, but does not run any statistical test to check if most deviating value is an outlier

`chisq.out.test()`: this function will check if extreme value is an outlier using hypothesis testing. The null hypothesis for the test is “H0: extreme value not an outlier”. Remember to look at p-value to make the decision whether to reject H0 or not.

`grubbs.test()`: this function will also check if extreme value is an outlier using hypothesis testing. The null hypothesis for the test is “H0: extreme value not an outlier”. Remember to look at p-value to make the decision whether to reject H0 or not.

`rm.outlier()`: if the result from the chi test tells you the extreme value is an outlier, then you can use this function to remove it or replace by sample mean or median.

When working with time series you cannot simply remove an outlier. Remember that in TSA we care about the time dependence structure, therefore eliminating observations is not an option. Instead we replace it with another value - preferably the local mean.



```
#Just find extreme value
outlier(inflow_random$inflow)
```


```
## [1] 4020.223
```

```
#Function chisq.out.test check if extreme value is outlier
chi_test <- chisq.out.test(inflow_random$inflow,var(inflow_random$inflow))
print(chi_test) #look at the p-value to find the decision
```

```
##
## chi-squared test for outlier
##
## data: inflow_random$inflow
## X-squared = 27.755, p-value = 1.377e-07
## alternative hypothesis: highest value 4020.22329300721 is an outlier
```

```
#If you need to remove outlier use rm.outlier()
inflow_random$inflow <- rm.outlier(inflow_random$inflow,fill=TRUE) #using fill equal true the value will
#Since we removed seasonality replacing with overall mean instead of local mean is acceptable
```

```
#Plot series again and look for more outliers
ggplot(inflow_random, aes(x=date, y=inflow)) +
  geom_line() +
  geom_abline(slope=0,intercept=3*sd_inflow,color="red") +
  geom_abline(slope=0,intercept=-3*sd_inflow,color="red")
```



Note we still have some outliers.

You can repeat the process until the next extreme value is not an outlier or write a loop as below.

```
summary(inflow_random$inflow)
```

```
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## -2300.860 -373.265  -23.874   -9.054   331.148  3734.378
```

```
#Writing a loop to remove all outliers
```

```
#Loop while to remove all outliers
```

```
pvalue <- 0 #just making sure we enter the while loop
```

```
aux_inflow <- inflow_random$inflow #Create a new vector for inflow_random just to make sure we don't l
```

```
nout <- 0 #keep track of number of outliers removed
```

```
while(pvalue < 0.05){ #the algorithm only enter the loop if the p-value
```

```
    #of first chi_test is less than 0.05 i.e. if there
```

```
    #is an outlier that needs to be removed
```

```
    out_test <- grubbs.test(aux_inflow,type=10)
```

```
    pvalue <- out_test$p.value #Update p-value every time we run the test for a new Aux_Y
```

```
    if(pvalue < 0.05){
```

```
        aux_inflow <- rm.outlier(aux_inflow,fill=TRUE) #replacing outliers
```

```
        nout <- nout+1
```

```
    }
```

```
}
```

```
cat("Number of outliers removed: ",nout,"\n")
```

```
## Number of outliers removed: 8
```

```
#Replaced original data with data without outliers
```

```
inflow_random$inflow <- aux_inflow
```


```
#Do the plots again
```

```
ggplot(inflow_random, aes(x=date, y=inflow)) +
```


```
    geom_line() +
```

```
    geom_abline(slope=0,intercept=3*sd_inflow,color="red") +
```

```
    geom_abline(slope=0,intercept=-3*sd_inflow,color="red")
```



```
ggplot(inflow_random, aes(y=inflow)) +  
  geom_boxplot()
```



*#Check the data*

```
summary(inflow_random$inflow)
```

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	-2300.86	-373.26	-28.45	-34.42	317.85	2631.74